

# StockMarketSuccess

*Dvir Blander and Corey Kozlovski*

*4/26/2019*

## R Markdown

### Goal and Description of Project

Our goal for this project is to be able to predict stocks by combining four methods of forecasting and weighing each one based on their individual successes in order to create one method that can accurately predict each stock individually. More importantly, we want to learn and understand new methods of analyzing data and being able to apply them in real world situations. Specifically, we will look at the stock of “Disney”, “Tesla”, “Coca Cola” and based on our analysis, we will determine the future trends of these stocks. We will be using Time Series, Neural Networks, Linear Regression, and lastly our own qualitative analysis of current events in order to come to an aggregate conclusion for each stock.

### Importing All the Packages and loading all the Libraries

```
library(quantmod)

## Loading required package: xts
## Loading required package: zoo
##
## Attaching package: 'zoo'
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
## Loading required package: TTR
## Version 0.4-0 included new data defaults. See ?getSymbols.
library(TimeWarp)
library(xts)
library(forecast)
library(zoo)
library(lubridate)

##
## Attaching package: 'lubridate'
## The following object is masked from 'package:base':
##
##   date
library(neuralnet)
library(BBmisc)

##
## Attaching package: 'BBmisc'
```

```
## The following object is masked from 'package:base':
##
## isFALSE
```

Getting the Tesla Data from the quantmod library and a plot of the closing values.

```
getSymbols(Symbols = "TSLA", auto.assign = TRUE)
```

```
## 'getSymbols' currently uses auto.assign=TRUE by default, but will
## use auto.assign=FALSE in 0.5-0. You will still be able to use
## 'loadSymbols' to automatically load data. getOption("getSymbols.env")
## and getOption("getSymbols.auto.assign") will still be checked for
## alternate defaults.
```

```
##
## This message is shown once per session and may be disabled by setting
## options("getSymbols.warning4.0"=FALSE). See ?getSymbols for details.
```

```
## [1] "TSLA"
```

```
plot(TSLA$TSLA.Close)
```



Getting the Coca Cola Data from the quantmod library and a plot of the closing values

```
getSymbols(Symbols = "KO", auto.assign = TRUE)
```

```
## [1] "KO"
```

```
plot(KO$KO.Close)
```



Getting the Disney Data from the quantmod library and a plot of the closing values.


```
getSymbols(Symbols = "DIS", auto.assign = TRUE)
```

```
## [1] "DIS"
```

```
plot(DIS$DIS.Close)
```

DIS\$DIS.Close

2007-01-03 / 2019-04-25



## Description for Current Event Analysis

Lastly, we will be using a more qualitative form of analysis through the means of current events. This analysis is very different than the other three, but we think it will bring very useful insight and context to our predictions from the other models. As this is a more general analysis, it will be most useful in predicting overall trends in the company's future in regards to real-life situations. Some drawbacks include being too reactionary to current events, as that would impact our predictions greatly, and furthermore this method isn't scientifically based as the other three are, but despite these shortcomings, we believe the context it provides will be essential to our predictions.

Current events for Tesla: When it comes to Tesla and News, these two terms almost come hand in hand. Be it from Elon Musk's Twitter antics, their full line-up of S-3-X-Y cars, and announcements of a new semi-truck and sports car, it can be very hard to use current events in predicting Tesla's future. They have some amazing developments, including battery recycling and releasing new cars recently that are more budget friendly, but also face some sharp downfalls. These include potential SEC violations that Musk faces due to his antics, a slowing down of car manufacturing for Q1 of 2019, and having the electric car as a whole not fully developed and realized hurts them as well. Overall, it is hard to predict the future of Tesla based on current events, but within the context of society and the growth of electric vehicles, it can be safe to assume Tesla will grow.

Current events for Coca Cola: The current events for Coca Cola are as follows. There is going to be advertising for Coca Cola in the newest Star Wars movie, which may bring significant revenue for the company. It is so a recession-proof business that has a healthy dividend, source: <https://www.fool.com/investing/2019/04/14/3-reasons-coca-cola-is-a-buy.aspx>. Additionally, it is a product that many people want even though people may try for healthier alternatives. Thus, expect continuous growth from Coca Cola.

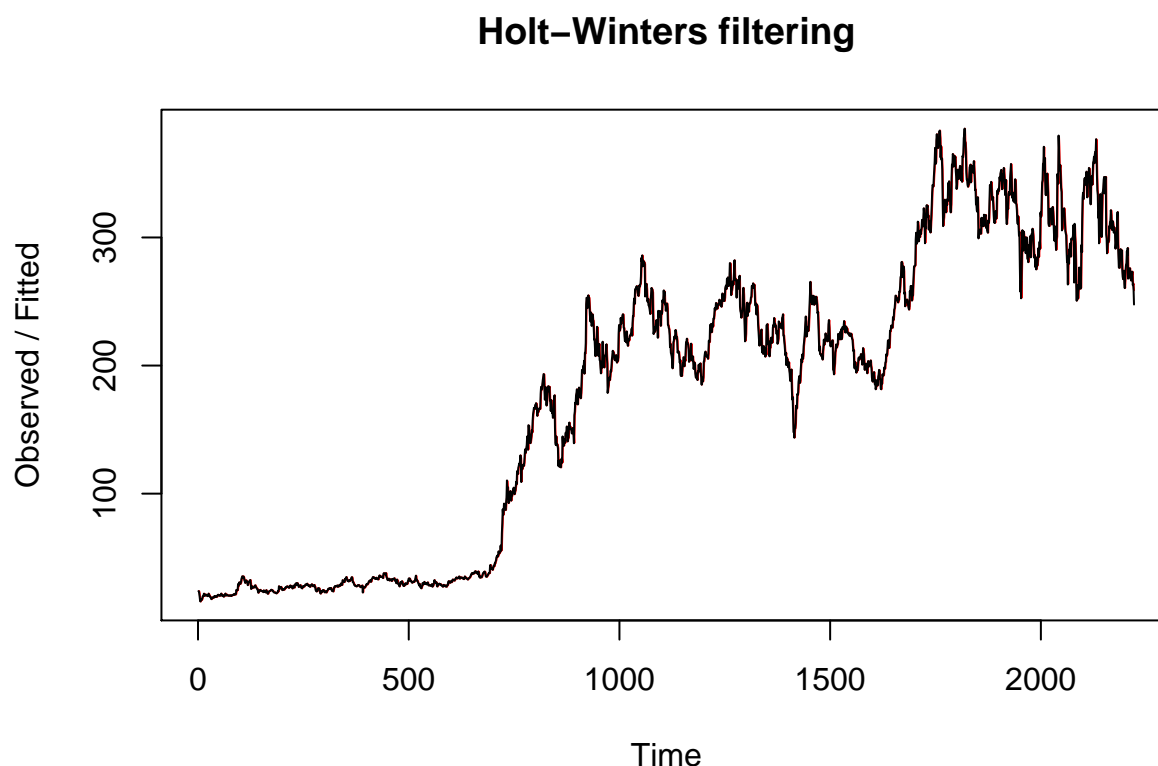
Current events for Disney: The current events for Disney are as follows. The company is releasing a streaming service called Disney Plus. It is going to be cheaper than Netflix but it is not expected to create much of a dent into Netflix's market share but, it should create some slow growth for the company. Source: <https://www.cnn.com/2019/04/15/what-a-1000-dollar-investment-in-disney-10-years-ago-would-be-worth-now.html>

## Description for Time Series

Starting with Time Series, more specifically through the Holt-Winters method, we understand that it is good at seeing long term trends overall and only relies on historical data, which is what we have through the means of the “quantmod” library in R. Some problems in relation to Holt-Winters Time Series is that historical data may not be particularly useful in forecasting future stocks and that it doesn’t take current events into account. To do this, we will be using the “forecast” library in conjunction with the native Holt-Winters Time Series method in R and analyzing the trend it predicts from the data.

Here is the Holt-Winters code for the Tesla dataset. First, we create the HoltWinters data set shown below. Here is a plot of the graph predicted by HoltWinters. This is the code that has HoltWinters predict what the future of the TSLA stock could be.

```
TSLA.pred <- HoltWinters(TSLA$TSLA.Close, beta = FALSE, gamma = FALSE)
plot(TSLA.pred)
```



```
TSLA.pred2 <- forecast::forecast.HoltWinters(TSLA.pred, h = 50)
TSLA.pred2
```

##	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
##	2222	247.8599	240.0431	255.6767	235.9052	259.8147
##	2223	247.8599	236.9188	258.8010	231.1269	264.5929
##	2224	247.8599	234.5065	261.2133	227.4376	268.2822
##	2225	247.8599	232.4677	263.2521	224.3196	271.4002
##	2226	247.8599	230.6691	265.0508	221.5688	274.1510
##	2227	247.8599	229.0416	266.6783	219.0797	276.6401
##	2228	247.8599	227.5440	268.1758	216.7894	278.9304
##	2229	247.8599	226.1495	269.5703	214.6567	281.0631
##	2230	247.8599	224.8393	270.8805	212.6530	283.0669
##	2231	247.8599	223.5998	272.1200	210.7573	284.9626
##	2232	247.8599	222.4206	273.2992	208.9539	286.7660

```

## 2233      247.8599 221.2937 274.4261 207.2304 288.4894
## 2234      247.8599 220.2127 275.5072 205.5771 290.1427
## 2235      247.8599 219.1724 276.5475 203.9861 291.7337
## 2236      247.8599 218.1685 277.5513 202.4508 293.2690
## 2237      247.8599 217.1975 278.5224 200.9657 294.7541
## 2238      247.8599 216.2562 279.4636 199.5263 296.1936
## 2239      247.8599 215.3423 280.3776 198.1285 297.5914
## 2240      247.8599 214.4533 281.2666 196.7689 298.9510
## 2241      247.8599 213.5873 282.1325 195.4445 300.2753
## 2242      247.8599 212.7428 282.9771 194.1529 301.5670
## 2243      247.8599 211.9180 283.8018 192.8915 302.8283
## 2244      247.8599 211.1118 284.6081 191.6585 304.0614
## 2245      247.8599 210.3228 285.3970 190.4519 305.2679
## 2246      247.8599 209.5502 286.1697 189.2702 306.4496
## 2247      247.8599 208.7928 286.9271 188.1119 307.6080
## 2248      247.8599 208.0498 287.6701 186.9755 308.7443
## 2249      247.8599 207.3204 288.3995 185.8600 309.8598
## 2250      247.8599 206.6039 289.1159 184.7643 310.9556
## 2251      247.8599 205.8996 289.8202 183.6872 312.0326
## 2252      247.8599 205.2070 290.5128 182.6279 313.0919
## 2253      247.8599 204.5254 291.1944 181.5856 314.1343
## 2254      247.8599 203.8544 291.8654 180.5593 315.1605
## 2255      247.8599 203.1935 292.5263 179.5485 316.1713
## 2256      247.8599 202.5422 293.1776 178.5525 317.1674
## 2257      247.8599 201.9002 293.8197 177.5705 318.1493
## 2258      247.8599 201.2670 294.4529 176.6021 319.1177
## 2259      247.8599 200.6422 295.0776 175.6467 320.0731
## 2260      247.8599 200.0257 295.6942 174.7037 321.0161
## 2261      247.8599 199.4169 296.3029 173.7728 321.9471
## 2262      247.8599 198.8158 296.9041 172.8534 322.8665
## 2263      247.8599 198.2219 297.4979 171.9451 323.7747
## 2264      247.8599 197.6350 298.0848 171.0476 324.6722
## 2265      247.8599 197.0550 298.6649 170.1604 325.5594
## 2266      247.8599 196.4814 299.2384 169.2833 326.4365
## 2267      247.8599 195.9142 299.8056 168.4158 327.3040
## 2268      247.8599 195.3531 300.3667 167.5577 328.1621
## 2269      247.8599 194.7980 300.9218 166.7087 329.0111
## 2270      247.8599 194.2486 301.4712 165.8685 329.8513
## 2271      247.8599 193.7048 302.0150 165.0368 330.6830


```

```

forecast::plot.forecast(TSLA.pred2, xlab = "Number of Days Recorded with Quantmod", ylab = "Stock Price")

```

## HoltWinters Tesla



Number of Days Recorded with Quantmod


This

shows the expected range in stock for the next 50 days.

Here is the Holt-Winters code for the Coca Cola dataset. First, we create the HoltWinters data set shown below. Here is a plot of the graph predicted by HoltWinters. This is the code that has HoltWinters predict what the future of the TSLA stock could be.

```
K0.pred <- HoltWinters(K0$K0.Close, beta = FALSE, gamma = FALSE)
plot(K0.pred)
```

## Holt-Winters filtering



```
KO.pred2 <- forecast::forecast.HoltWinters(KO.pred, h = 50)
KO.pred2
```


##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 3100	47.84671	47.36704	48.32638	47.11312	48.58030
## 3101	47.84671	47.18348	48.50994	46.83238	48.86103
## 3102	47.84671	47.04069	48.65273	46.61400	49.07941
## 3103	47.84671	46.91963	48.77378	46.42887	49.26454
## 3104	47.84671	46.81266	48.88076	46.26526	49.42815
## 3105	47.84671	46.71576	48.97766	46.11706	49.57635
## 3106	47.84671	46.62652	49.06689	45.98060	49.71282
## 3107	47.84671	46.54339	49.15003	45.85345	49.83996
## 3108	47.84671	46.46525	49.22817	45.73394	49.95947
## 3109	47.84671	46.39129	49.30212	45.62084	50.07257
## 3110	47.84671	46.32092	49.37250	45.51322	50.18020
## 3111	47.84671	46.25365	49.43976	45.41034	50.28307
## 3112	47.84671	46.18911	49.50430	45.31164	50.38178
## 3113	47.84671	46.12700	49.56642	45.21663	50.47678
## 3114	47.84671	46.06704	49.62637	45.12495	50.56847
## 3115	47.84671	46.00905	49.68437	45.03625	50.65717
## 3116	47.84671	45.95282	49.74059	44.95026	50.74315
## 3117	47.84671	45.89822	49.79519	44.86676	50.82666
## 3118	47.84671	45.84511	49.84830	44.78553	50.90788
## 3119	47.84671	45.79337	49.90004	44.70640	50.98701
## 3120	47.84671	45.74291	49.95051	44.62922	51.06419
## 3121	47.84671	45.69362	49.99979	44.55385	51.13957
## 3122	47.84671	45.64544	50.04797	44.48016	51.21325
## 3123	47.84671	45.59830	50.09512	44.40806	51.28536
## 3124	47.84671	45.55212	50.14130	44.33743	51.35598
## 3125	47.84671	45.50685	50.18657	44.26820	51.42522



```
## 3126    47.84671 45.46244 50.23098 44.20028 51.49313
## 3127    47.84671 45.41884 50.27458 44.13361 51.55981
## 3128    47.84671 45.37601 50.31740 44.06811 51.62531
## 3129    47.84671 45.33392 50.35950 44.00372 51.68969
## 3130    47.84671 45.29251 50.40091 43.94040 51.75302
## 3131    47.84671 45.25177 50.44165 43.87809 51.81533
## 3132    47.84671 45.21165 50.48176 43.81674 51.87668
## 3133    47.84671 45.17214 50.52127 43.75631 51.93710
## 3134    47.84671 45.13320 50.56021 43.69676 51.99665
## 3135    47.84671 45.09482 50.59860 43.63806 52.05536
## 3136    47.84671 45.05696 50.63646 43.58016 52.11326
## 3137    47.84671 45.01961 50.67381 43.52303 52.17038
## 3138    47.84671 44.98275 50.71067 43.46666 52.22676
## 3139    47.84671 44.94635 50.74707 43.41099 52.28242
## 3140    47.84671 44.91041 50.78301 43.35602 52.33739
## 3141    47.84671 44.87490 50.81852 43.30171 52.39170
## 3142    47.84671 44.83981 50.85361 43.24805 52.44537
## 3143    47.84671 44.80512 50.88830 43.19500 52.49841
## 3144    47.84671 44.77083 50.92259 43.14255 52.55086
## 3145    47.84671 44.73691 50.95651 43.09068 52.60273
## 3146    47.84671 44.70336 50.99006 43.03937 52.65404
## 3147    47.84671 44.67016 51.02325 42.98860 52.70481
## 3148    47.84671 44.63731 51.05610 42.93836 52.75506
## 3149    47.84671 44.60479 51.08862 42.88862 52.80479
```

```
forecast::plot.forecast(KO.pred2, xlab = "Number of Days Recorded with Quantmod", ylab = "Stock Price")
```

## HoltWinters Coca Cola




This shows the expected range in stock for the next 50 days.

Here is the Holt-Winters code for the Disney dataset. First, we create the HoltWinters data set shown below. Here is a plot of the graph predicted by HoltWinters. This is the code that has HoltWinters predict what the

future of the TSLA stock could be.

```
DIS.pred <- HoltWinters(DIS$DIS.Close, beta = FALSE, gamma = FALSE)
plot(DIS.pred)
```

## Holt-Winters filtering



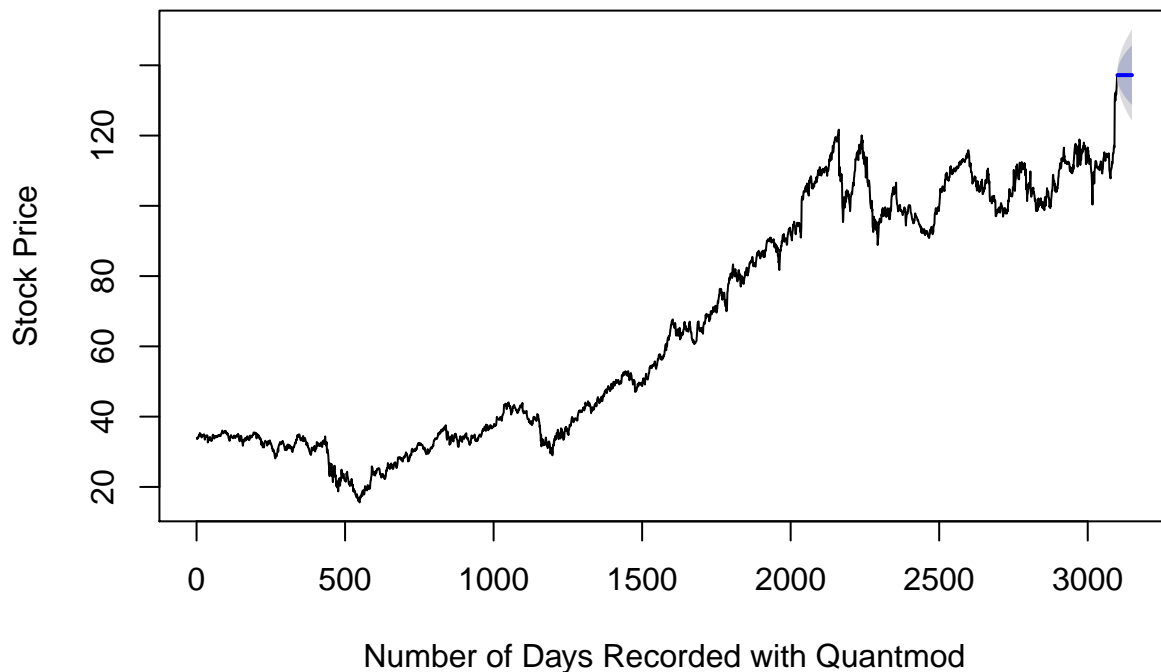
```
DIS.pred2 <- forecast::forecast.HoltWinters(DIS.pred, h = 50)
DIS.pred2
```

##	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
##	3100	137.2035	135.9835	138.4235	135.3377	139.0694
##	3101	137.2035	135.4926	138.9144	134.5869	139.8201
##	3102	137.2035	135.1140	139.2930	134.0079	140.3991
##	3103	137.2035	134.7942	139.6128	133.5188	140.8882
##	3104	137.2035	134.5121	139.8949	133.0874	141.3196
##	3105	137.2035	134.2569	140.1501	132.6971	141.7099
##	3106	137.2035	134.0222	140.3849	132.3380	142.0690
##	3107	137.2035	133.8035	140.6035	132.0037	142.4033
##	3108	137.2035	133.5982	140.8089	131.6896	142.7174
##	3109	137.2035	133.4039	141.0032	131.3924	143.0146
##	3110	137.2035	133.2190	141.1880	131.1098	143.2973
##	3111	137.2035	133.0424	141.3646	130.8396	143.5674
##	3112	137.2035	132.8730	141.5341	130.5805	143.8265
##	3113	137.2035	132.7099	141.6971	130.3311	144.0759
##	3114	137.2035	132.5526	141.8544	130.0905	144.3165
##	3115	137.2035	132.4004	142.0066	129.8578	144.5493
##	3116	137.2035	132.2529	142.1541	129.6322	144.7748
##	3117	137.2035	132.1096	142.2974	129.4131	144.9939
##	3118	137.2035	131.9703	142.4367	129.2000	145.2070
##	3119	137.2035	131.8346	142.5724	128.9925	145.4145
##	3120	137.2035	131.7023	142.7048	128.7901	145.6170
##	3121	137.2035	131.5730	142.8340	128.5924	145.8146

```
## 3122      137.2035 131.4467 142.9604 128.3992 146.0079
## 3123      137.2035 131.3230 143.0840 128.2101 146.1969
## 3124      137.2035 131.2019 143.2051 128.0249 146.3821
## 3125      137.2035 131.0832 143.3238 127.8434 146.5637
## 3126      137.2035 130.9668 143.4402 127.6653 146.7417
## 3127      137.2035 130.8525 143.5545 127.4905 146.9165
## 3128      137.2035 130.7402 143.6668 127.3188 147.0882
## 3129      137.2035 130.6299 143.7772 127.1500 147.2570
## 3130      137.2035 130.5213 143.8857 126.9840 147.4230
## 3131      137.2035 130.4145 143.9925 126.8207 147.5864
## 3132      137.2035 130.3094 144.0976 126.6598 147.7472
## 3133      137.2035 130.2058 144.2012 126.5014 147.9056
## 3134      137.2035 130.1037 144.3033 126.3454 148.0617
## 3135      137.2035 130.0031 144.4039 126.1915 148.2155
## 3136      137.2035 129.9039 144.5031 126.0397 148.3673
## 3137      137.2035 129.8060 144.6010 125.8900 148.5170
## 3138      137.2035 129.7094 144.6976 125.7423 148.6648
## 3139      137.2035 129.6140 144.7930 125.5964 148.8106
## 3140      137.2035 129.5198 144.8872 125.4523 148.9547
## 3141      137.2035 129.4267 144.9803 125.3100 149.0970
## 3142      137.2035 129.3348 145.0722 125.1693 149.2377
## 3143      137.2035 129.2439 145.1631 125.0303 149.3767
## 3144      137.2035 129.1540 145.2530 124.8929 149.5141
## 3145      137.2035 129.0651 145.3419 124.7569 149.6501
## 3146      137.2035 128.9772 145.4298 124.6225 149.7845
## 3147      137.2035 128.8902 145.5168 124.4894 149.9176
## 3148      137.2035 128.8041 145.6029 124.3578 150.0492
## 3149      137.2035 128.7189 145.6881 124.2275 150.1796
```

```
forecast::plot.forecast(DIS.pred2, xlab = "Number of Days Recorded with Quantmod", ylab = "Stock Price")
```

## HoltWinters Disney



This

shows the expected range in stock for the next 50 days.

## Description for Neural Networks

Moving forward with Neural Networks, we will be using the “neuralnet”, “xts”, “zoo”, “lubridate”, “BBmisc” libraries in R. We understand that it is beneficial for larger datasets and that it can be more efficient in terms of time to be run. However, our datasets are not in the size of millions and it is difficult to build and understand a neural network algorithm. Thus, it is to be expected that this method will be weighed less than other methods.

Here is the code for Tesla. We first start by creating a new dataframe, as the values we are using to learn are different than the other methods. We also have to filter the dataset, as seen in the for loop.

```
TSLA.df <- data.frame("Open" = TSLA$TSLA.Open, "Close" = TSLA$TSLA.Close, "Year"=as.numeric(format(index(TSLA), "%Y")),
  "Month"=as.numeric(format(index(TSLA), "%m")), "Day"=as.numeric(format(index(TSLA), "%d")),
  "Is_Month_Start" = FALSE, "Is_Month_End" = FALSE, "Is_Year_End" = FALSE,
  "Is_Year_Start" = FALSE, "Is_Quarter_Start" = FALSE, "Is_Quarter_End" = FALSE,
  "Quarters"=quarters(index(TSLA)))

for (i in 1:nrow(TSLA.df)){

  date <- as.POSIXlt(paste(TSLA.df[i,]$Year, "-", TSLA.df[i,]$Month, "-", TSLA.df[i,]$Day, sep=""))

  #Determine if it is beginning/end of the month and change dataset accordingly
  if (TSLA.df[i,]$Day <= 7){ #7 is the first ~25% of the month
    TSLA.df[i,]$Is_Month_Start = TRUE
  } else if (TSLA.df[i,]$Day > 24){ #24 is the last ~25% of the month
    TSLA.df[i,]$Is_Month_End = TRUE
  }

  #Determine if it is the beginning/end of the year and change dataset accordingly
  if (TSLA.df[i,]$Quarters == 'Q1'){ #Q1 is the first ~25% of the year
    TSLA.df[i,]$Is_Year_Start = TRUE
  } else if (TSLA.df[i,]$Quarters == 'Q4'){ #Q4 is the last ~25% of the year
    TSLA.df[i,]$Is_Year_End = TRUE
  }

  #Determine if it is the beginning of the quarter and change dataset accordingly
  if (yday(date) <= 22){ #days 1-22 is the first ~25% of Q1
    TSLA.df[i,]$Is_Quarter_Start = TRUE
  } else if (yday(date) > 91 & yday(date) <= 113){ #days 91-113 is the first ~25% of Q2
    TSLA.df[i,]$Is_Quarter_Start = TRUE
  } else if (yday(date) > 182 & yday(date) <= 204){ #days 182-204 is the first ~25% of Q3
    TSLA.df[i,]$Is_Quarter_Start = TRUE
  } else if (yday(date) > 273 & yday(date) <= 296){ #days 273-296 is the first ~25% of Q4
    TSLA.df[i,]$Is_Quarter_Start = TRUE
  }

  #Determine if it is the end of the quarter and change dataset accordingly
  if (yday(date) <= 91 & yday(date) >= 69){ #days 69-91 is the last ~25% of Q1
    TSLA.df[i,]$Is_Quarter_End = TRUE
  } else if (yday(date) <= 182 & yday(date) >= 160){ #days 160-182 is the last ~25% of Q2
    TSLA.df[i,]$Is_Quarter_End = TRUE
  } else if (yday(date) <= 273 & yday(date) >= 251){ #days 251-273 is the last ~25% of Q3
    TSLA.df[i,]$Is_Quarter_End = TRUE
  }
}
```

```

    TSLA.df[i,]$Is_Quarter_End = TRUE
  } else if (yday(date) <= 366 & yday(date) >= 343){ #days 343-365 is the last ~25% of Q4
    TSLA.df[i,]$Is_Quarter_End = TRUE
  }
}

TSLA.df <- TSLA.df[,1:11] #Remove the last column, as it was only used to clean data

#change true/false values into 1/0s to prepare for normalization.
cols <- sapply(TSLA.df, is.logical)
TSLA.df[,cols] <- lapply(TSLA.df[,cols], as.numeric)

```

To use neural networks, we have to normalize the data set. We used the BBmisc library to help us with this.

```
TSLA.df <- BBmisc::normalize(TSLA.df)
```

Now, we make a testing and training set to evaluate the performance of neural networks. We decided to use 75% of the original data set.

```

TSLA.train <- TSLA.df[1:(.75 * nrow(TSLA.df)),]
TSLA.test <- TSLA.df[(.75 * nrow(TSLA.df)):nrow(TSLA.df),]

```

We now create the neural network and display it.

```

TSLA.NN <- neuralnet(TSLA.Close ~ TSLA.Open + Year + Month + Day + Is_Month_Start + Is_Month_End + Is_Y
                      Is_Year_Start + Is_Quarter_Start + Is_Quarter_End
                      , data = TSLA.train, linear.output=TRUE, hidden=c(5,3,3))

plot(TSLA.NN)

```

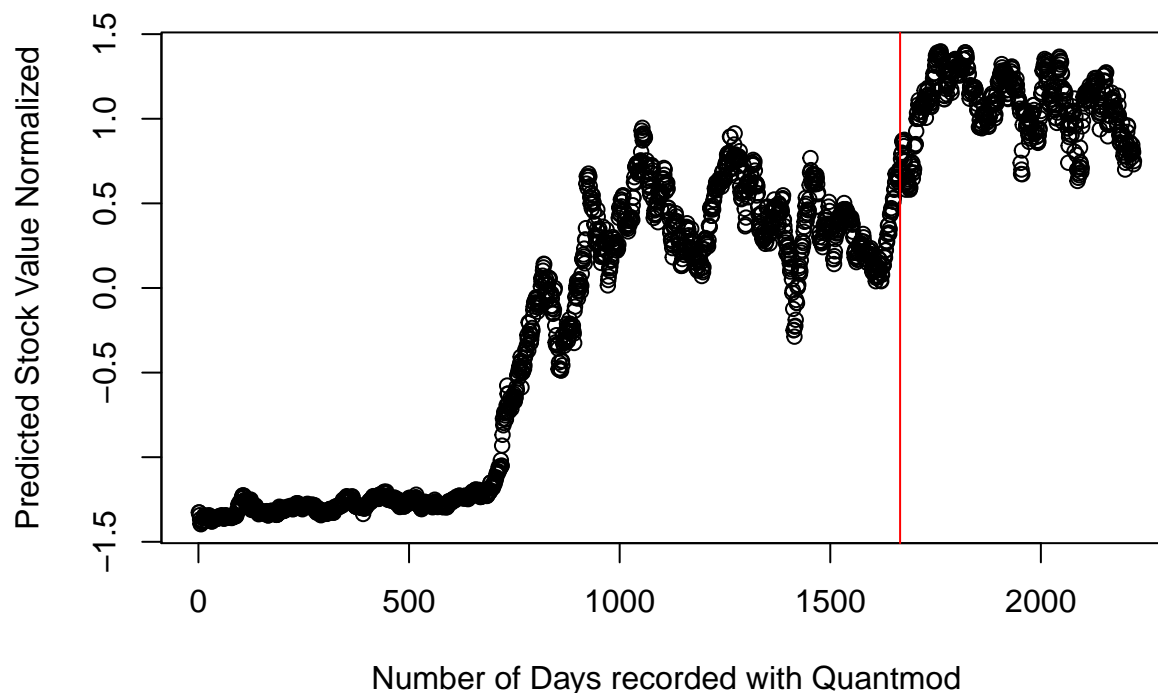
Using the created neural network, we predict the values and compare them to the actual values. The red line indicates where the training data ends and where the prediction and actual values, respectively, begin.

```

TSLA.pred <- predict(TSLA.NN, TSLA.test)
TSLA.pred.total <- c(TSLA.train$TSLA.Close, TSLA.pred)
plot(TSLA.pred.total, xlab="Number of Days recorded with Quantmod", ylab="Predicted Stock Value Normalized",
     abline(v = (.75 * nrow(TSLA.df)), col='red'))

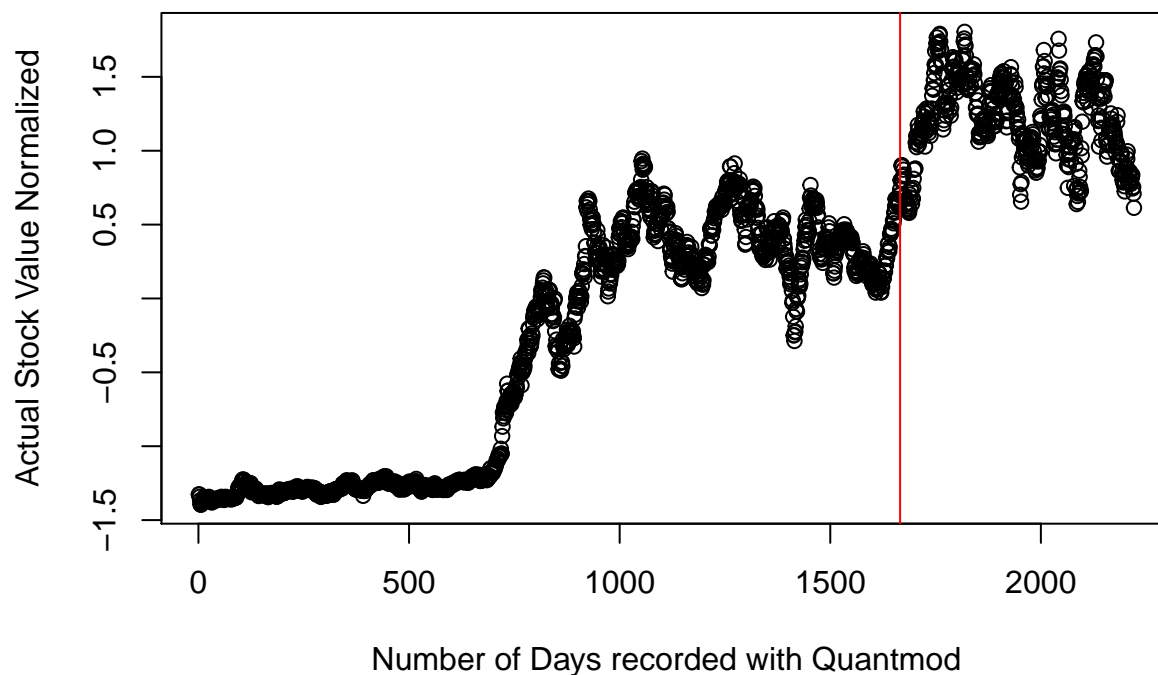
```

## Normalized Stock Predictions of TSLA



```
plot(TSLA.df$TSLA.Close, xlab="Number of Days recorded with Quantmod", ylab="Actual Stock Value Normalized", col='green')  
abline(v = (.75 * nrow(TSLA.df)), col='red')
```

## Actual Stock Values of TSLA



Comparing both graphs, we can see that our predictor did a pretty poor job of predicting the values for TSLA. It predicted much less volatility than what actually happened, and did not predict it to go nearly as high as it did based on the normalized scale. The prediction maxed out at just above 1, but the actual values went

more towards 1.5. Other than that, it did manage to predict that there would be 3 large “dips”, as both the predicted graph and the actual graph show it. I would not recommend using this method, with the parameters we chose, to predict, but it did do well in seeing dips.

Here is the code for Coca Cola. We first start by creating a new dataframe, as the values we are using to learn are different than the other methods. We also have to filter the dataset, as seen in the for loop.

```
KO.df <- data.frame("Open" = KO$KO.Open, "Close" = KO$KO.Close, "Year"=as.numeric(format(index(KO), "%Y")),
                    "Month"=as.numeric(format(index(KO), "%m")), "Day"=as.numeric(format(index(KO), "%d")),
                    "Is_Month_Start" = FALSE, "Is_Month_End" = FALSE, "Is_Year_End" = FALSE,
                    "Is_Year_Start" = FALSE, "Is_Quarter_Start" = FALSE, "Is_Quarter_End" = FALSE,
                    "Quarters"=quarters(index(KO)))

for (i in 1:nrow(KO.df)){

  date <- as.POSIXlt(paste(KO.df[i,]$Year, "-", KO.df[i,]$Month, "-", KO.df[i,]$Day, sep=""))

  #Determine if it is beginning/end of the month and change dataset accordingly
  if (KO.df[i,]$Day <= 7){ #7 is the first ~25% of the month
    KO.df[i,]$Is_Month_Start = TRUE
  } else if (KO.df[i,]$Day > 24){ #24 is the last ~25% of the month
    KO.df[i,]$Is_Month_End = TRUE
  }

  #Determine if it is the beginning/end of the year and change dataset accordingly
  if (KO.df[i,]$Quarters == 'Q1'){ #Q1 is the first ~25% of the year
    KO.df[i,]$Is_Year_Start = TRUE
  } else if (KO.df[i,]$Quarters == 'Q4'){ #Q4 is the last ~25% of the year
    KO.df[i,]$Is_Year_End = TRUE
  }

  #Determine if it is the beginning of the quarter and change dataset accordingly
  if (yday(date) <= 22){ #days 1-22 is the first ~25% of Q1
    KO.df[i,]$Is_Quarter_Start = TRUE
  } else if (yday(date) > 91 & yday(date) <= 113){ #days 91-113 is the first ~25% of Q2
    KO.df[i,]$Is_Quarter_Start = TRUE
  } else if (yday(date) > 182 & yday(date) <=204){ #days 182-204 is the first ~25% of Q3
    KO.df[i,]$Is_Quarter_Start = TRUE
  } else if (yday(date) > 273 & yday(date) <=296){ #days 273-296 is the first ~25% of Q4
    KO.df[i,]$Is_Quarter_Start = TRUE
  }

  #Determine if it is the end of the quarter and change dataset accordingly
  if (yday(date) <= 91 & yday(date) >= 69){ #days 69-91 is the last ~25% of Q1
    KO.df[i,]$Is_Quarter_End = TRUE
  } else if (yday(date) <= 182 & yday(date) >= 160){ #days 160-182 is the last ~25% of Q2
    KO.df[i,]$Is_Quarter_End = TRUE
  } else if (yday(date) <= 273 & yday(date) >= 251){ #days 251-273 is the last ~25% of Q3
    KO.df[i,]$Is_Quarter_End = TRUE
  } else if (yday(date) <= 366 & yday(date) >= 343){ #days 343-365 is the last ~25% of Q4
    KO.df[i,]$Is_Quarter_End = TRUE
  }
}

KO.df <- KO.df[,1:11]#Remove the last column, it was only used to clean data
```

```
#change true/false values into 1/0s
cols <- sapply(KO.df, is.logical)
KO.df[,cols] <- lapply(KO.df[,cols], as.numeric)
```

To use neural networks, we have to normalize the data set. We used the BBmisc library to help us with this.

```
KO.df <- BBmisc::normalize(KO.df)
```

Now, we make a testing and training set to evaluate the performance of neural networks. We decided to use 75% of the original data set.

```
KO.train <- KO.df[1:(.75 * nrow(KO.df)),] #75% of data set
KO.test <- KO.df[(.75 * nrow(KO.df)):nrow(KO.df),]
```


We now create the neural network and display it.

```
KO.NN <- neuralnet(KO.Close ~ KO.Open + Year + Month + Day + Is_Month_Start + Is_Month_End + Is_Year_End +
                    Is_Year_Start + Is_Quarter_Start + Is_Quarter_End
                    , data = KO.train, linear.output=TRUE, hidden=c(5,3,3))
plot(KO.NN)
```

Using the created neural network, we predict the values and compare them to the actual values. The red line indicates where the training data ends and where the prediction and actual values, respectively, begin.

```
KO.pred <- predict(KO.NN, KO.test)
KO.pred.total <- c(KO.train$KO.Close, KO.pred)
plot(KO.pred.total, xlab="Number of Days recorded with Quantmod", ylab="Predicted Stock Value Normalized",
      abline(v = (.75 * nrow(KO.df)), col='red'))
```


## Normalized Stock Predictions of KO



```
plot(KO.df$KO.Close, xlab="Number of Days recorded with Quantmod", ylab="Actual Stock Value Normalized",
      abline(v = (.75 * nrow(KO.df)), col='red'))
```



## Actual Stock Values of KO



### Number of Days recorded with Quantmod

Surprisingly, this model worked very well for Coca Cola. It managed to predict the 3 peaks fairly accurately and both dips as well. It is nearly exact overall but the only significant difference would be that the neural network was significantly less aggressive when predicting peaks. The actual values ended up going a bit higher than predicted, but they were very close. For this data set, I would recommend using this method.

Here is the code for Disney. We first start by creating a new dataframe, as the values we are using to learn are different than the other methods. We also have to filter the dataset, as seen in the for loop.

```
DIS.df <- data.frame("Open" = DIS$DIS.Open, "Close" = DIS$DIS.Close, "Year"=as.numeric(format(index(DIS),
"Month"=as.numeric(format(index(DIS), "%m")), "Day"=as.numeric(format(index(DIS),
"Is_Month_Start" = FALSE, "Is_Month_End" = FALSE, "Is_Year_End" = FALSE,
"Is_Year_Start" = FALSE, "Is_Quarter_Start" = FALSE, "Is_Quarter_End" = FALSE,
"Quarters"=quarters(index(DIS)))

for (i in 1:nrow(DIS.df)){

  date <- as.POSIXlt(paste(DIS.df[i,]$Year, "-", DIS.df[i,]$Month, "-", DIS.df[i,]$Day, sep=""))

  #Determine if it is beginning/end of the month and change dataset accordingly
  if (DIS.df[i,]$Day <= 7){ #7 is the first ~25% of the month
    DIS.df[i,]$Is_Month_Start = TRUE
  } else if (DIS.df[i,]$Day > 24){ #24 is the last ~25% of the month
    DIS.df[i,]$Is_Month_End = TRUE
  }

  #Determine if it is the beginning/end of the year and change dataset accordingly
  if (DIS.df[i,]$Quarters == 'Q1'){ #Q1 is the first ~25% of the year
    DIS.df[i,]$Is_Year_Start = TRUE
  } else if (DIS.df[i,]$Quarters == 'Q4'){ #Q4 is the last ~25% of the year
    DIS.df[i,]$Is_Year_End = TRUE
  }
}
```

```

#Determine if it is the beginning of the quarter and change dataset accordingly
if (yday(date) <= 22){ #days 1-22 is the first ~25% of Q1
  DIS.df[i,]$Is_Quarter_Start = TRUE
} else if (yday(date) > 91 & yday(date) <= 113){ #days 91-113 is the first ~25% of Q2
  DIS.df[i,]$Is_Quarter_Start = TRUE
} else if (yday(date) > 182 & yday(date) <=204){ #days 182-204 is the first ~25% of Q3
  DIS.df[i,]$Is_Quarter_Start = TRUE
} else if (yday(date) > 273 & yday(date) <=296){ #days 273-296 is the first ~25% of Q4
  DIS.df[i,]$Is_Quarter_Start = TRUE
}

#Determine if it is the end of the quarter and change dataset accordingly
if (yday(date) <= 91 & yday(date) >= 69){ #days 69-91 is the last ~25% of Q1
  DIS.df[i,]$Is_Quarter_End = TRUE
} else if (yday(date) <= 182 & yday(date) >= 160){ #days 160-182 is the last ~25% of Q2
  DIS.df[i,]$Is_Quarter_End = TRUE
} else if (yday(date) <= 273 & yday(date) >= 251){ #days 251-273 is the last ~25% of Q3
  DIS.df[i,]$Is_Quarter_End = TRUE
} else if (yday(date) <= 366 & yday(date) >= 343){ #days 343-365 is the last ~25% of Q4
  DIS.df[i,]$Is_Quarter_End = TRUE
}
}

DIS.df <- DIS.df[,1:11]#Remove the last column, it was only used to clean data

#change true/false values into 1/0s
cols <- sapply(DIS.df, is.logical)
DIS.df[,cols] <- lapply(DIS.df[,cols], as.numeric)

```

To use neural networks, we have to normalize the data set. We used the BBmisc library to help us with this.

```
DIS.df <- BBmisc::normalize(DIS.df)
```

Now, we make a testing and training set to evaluate the performance of neural networks. We decided to use 75% of the original data set.

```
DIS.train <- DIS.df[1:(.75 * nrow(DIS.df)),] #75% of data set
DIS.test <- DIS.df[(.75 * nrow(DIS.df)):nrow(DIS.df),]
```

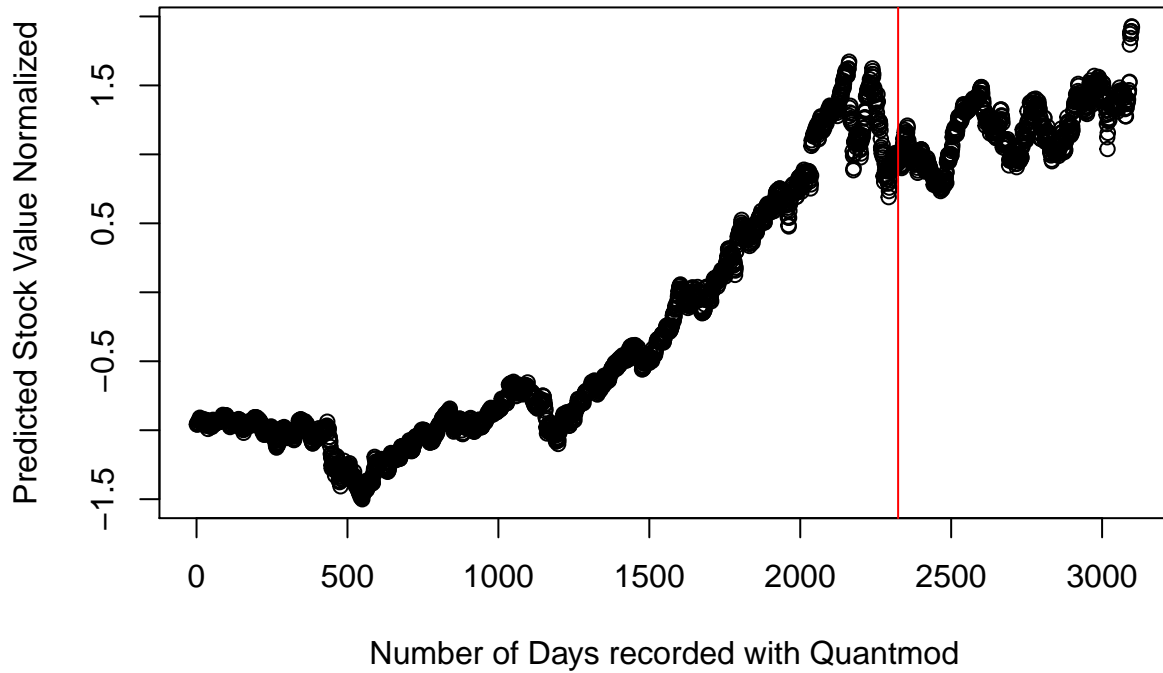
We now create the neural network and display it.

```
DIS.NN <- neuralnet(DIS.Close ~ DIS.Open + Year + Month + Day + Is_Month_Start + Is_Month_End + Is_Year_Start + Is_Year_End + Is_Quarter_Start + Is_Quarter_End
, data = DIS.train, linear.output=TRUE, hidden=c(5,3,3))
plot(DIS.NN)
```

Using the created neural network, we predict the values and compare them to the actual values. The red line indicates where the training data ends and where the prediction and actual values, respectively, begin.


```
DIS.pred <- predict(DIS.NN, DIS.test)
DIS.pred.total <- c(DIS.train$DIS.Close, DIS.pred)
plot(DIS.pred.total, xlab="Number of Days recorded with Quantmod", ylab="Predicted Stock Value Normalized",
abline(v = (.75 * nrow(DIS.df)), col='red')
```

## Normalized Stock Predictions of DIS



```
plot(DIS.df$DIS.Close, xlab="Number of Days recorded with Quantmod", ylab="Actual Stock Value Normalized", col='red')
abline(v = (.75 * nrow(DIS.df)), col='red')
```

## Actual Stock Values of DIS



Again, this model performed exceedingly well! It was able to predict 3 peaks, 3 dips, and even manage to predict, towards the very end, the start of a peak. As with the last one, Neural Networks seem to be less aggressive with the predictions overall, as the actual data set reached and exceeded 2, but the predictions didn't quite


get there. I would again recommend using this model for our overall stock market prediction.

## Description for Linear Regression

Onward, we will be looking into the Linear Regression method of predicting the future values of our chosen stocks. The method in which this will be executed is by choosing the most recent trend in the data by looking at the plot for each graph and finding a linear regression of the graphs to make predictions. Some benefits for this method is that it will use all of the data and fit the data in the best possible way and the model will also favor newer data rather than older data. Moreover, a property in Calculus is that if one restricts the domain in any graph enough, the graph will show up as linear. This is why Linear Regression is a valid form of analyzing the data. However, there is a risk in excluding necessary data to create the trend of the graph. Additionally, this is a pretty simplistic way of predicting future stocks. Thus, this may be weighed less than other options. The Timewarp library was used in addition to the quantmod library.

Here are the steps for Linear Regression in the Tesla dataset. First we restrict the domain of the dataset to only view the most recent trend of the data. We next use Chartseries with subset as the main argument. We then use a linear model on the same range of the chart in addition to plotting the chart.

```
subset <- TSLA[1700:nrow(TSLA),]  
chartSeries(subset, TA = NULL, theme = "white", up.col = "green", dn.col = "red")  
indices = 1:nrow(subset)  
model=lm(TSLA.Close~indices,data=subset)  
abline(model$coefficients[1],model$coefficients[2])
```




Based on the graph above, Tesla seems to be a stock that is decreasing in value.

Here are the steps for Linear Regression in the Coca Cola dataset. First we restrict the domain of the dataset

to only view the most recent trend of the data. We next use `Chartseries` with `subset` as the main argument. We then use a linear model on the same range of the chart in addition to plotting the chart.


```
subset = KO[1500:nrow(KO),]
chartSeries(subset, TA = NULL, theme = "white", up.col = "green", dn.col = "red")
indices = 1:nrow(subset)
model=lm(KO.Close~indices,data=subset)
abline(model$coefficients[1],model$coefficients[2])
```



Based on the graph above, Coca Cola seems to be a stock that is increasing in value.

Here are the steps for Linear Regression in the Disney dataset. First we restrict the domain of the dataset to only view the most recent trend of the data. We next use `Chartseries` with `subset` as the main argument. We then use a linear model on the same range of the chart in addition to plotting the chart.

```
subset <- DIS[2000:nrow(DIS),]
chartSeries(subset, TA = NULL, theme = "white", up.col = "green", dn.col = "red")
indices = 1:nrow(subset)
model=lm(DIS.Close~indices,data=subset)
abline(model$coefficients[1],model$coefficients[2])
```



Based on the graph above, Disney seems to be a stock that is increasing in value.

## Conclusion

Based on the pros and cons of each method of forecasting, we are going to weigh each model differently, placing higher weight on Time-Series and Current Events, and lower weights on Neural Networks and Linear Regression. The different weights for each model will be based on how we expect each model to contribute to our conclusion. Since our pros and cons for the Time Series and the Current Events has led us to believe more in those methods, they will have higher contribution than Neural Networks and Linear Regression. If we learn that our method is successful, we think it could be very useful in predicting the outcome of future stocks for a variety of other companies as well.

In terms of Tesla, we predicted that the current events would have an increase for the stock value. The Holt-winters analysis predicted a large range of values for the future, which creates a lot of uncertainty. The Linear regression analysis showed that the Tesla stock tended to decrease. The Neural Network analysis did not really fit the Tesla data and we wouldn't feel safe using this, in a different context, to predict Tesla's future. Thus, we are not sure what will happen with the Tesla stock in the future due to the conflicting analysis' for Linear Regression and Current Events.

For Coca Cola, we found that in almost every test that it would end up growing. For Current Events, we found that everything that is happening for that company overall benefits them and Holt Winters Time Series we found that the range of values isn't nearly as volatile as Tesla's. These two tests we valued higher than the other two, but despite that, we found that Linear Regression also increased alongside having Neural Networks be very accurate in predictions. With that knowledge, if we could apply Neural Networks in a different context and predict it, we believe it would also see an increase. Overall, all of our tests tend to lead towards Coca Cola only growing from here.

Finally, the Disney stock should grow slowly based on the following results. The Current Event analysis had the stock increasing, the Linear Regression indicated in an increasing stock value, and the Holt-Winters had a smaller range of possible values for Disney. The Neural Networks was also very accurate in its analysis. With that knowledge, if we could apply Neural Networks in a different context and predict it, we believe it would also see an increase.

From this project, we noticed a few trends that were apparent. First of all, if Holt Winters performed well and had not that much uncertainty, we found that Neural Networks would perform very well. If it did, Neural Networks struggled to perform, as seen in the tests ran on Tesla. In addition, we found that for all 3 companies, current events would always lead to us believing in an increase to the stock value. This could be very dependent on our biases, in the sense of how we value each event, in addition to news source articles that we received the information from. Overall, it was very interesting applying and learning new ways to predict real-life situations.